Response to the Reviewers

Paper ID 280, "Cost and Safety of Label Acquisition Strategies for MALDI-TOF Antimicrobial Resistance Prediction"
7th International Conference on Innovative Computing (ICIC 2026)

We thank both reviewers and the editor. Every comment was acted on. The points that needed new computation were recomputed from the stored run records, and one new experiment (1800 further runs) was carried out for Reviewer 2, point 2. To keep the paper within 6 pages, the submitted Table I (position of this study) and Fig. 4 (SHAP dot matrix) were removed, since Section II-C and Table III carry the same content, and a few sentences that repeated others were deleted. Section and table numbers below refer to the camera ready version.

Reviewer 2

1. Verify the Wilcoxon and Holm adjusted p values.

The reviewer and the editor are correct. The submitted text said that the Holm procedure was applied across the whole family of comparisons. Our analysis code applied it within each species and antibiotic pair, which is a family of 15 comparisons (3 rules at 5 budgets). The reported value of 0.029 is 15 times 0.00195, the smallest two sided value an exact Wilcoxon signed rank test can return on 10 seeds, so the number was produced by the test that ran and the description of that test was wrong.

All tests were recomputed from the per run records. Section IV-C now states the settings (exact two sided Wilcoxon signed rank test on the paired difference in test AUROC, 10 matched seeds per comparison in the temporal setting and 5 in the cross site setting, Holm correction both within each pair and across all 45 temporal comparisons) and the smallest attainable values (0.002 uncorrected, 0.029 within a pair, 0.088 across all 45, and 0.0625 on 5 seeds). Section V-B reports both corrections. Within a pair, 2 comparisons remain below 0.05 and both favour random selection. Across all 45 the smallest value is 0.088, so no comparison is significant in either direction. The abstract and the conclusion no longer state that 2 rules were worse. As requested, effect sizes are reported with 95% confidence intervals. The new Table I gives the mean paired difference in AUROC and its interval for all 45 temporal comparisons, and Section V-B gives the two largest losses (uncertainty sampling on *K. pneumoniae* at 800 labels, 0.029, 95% interval 0.014 to 0.045, and diversity sampling on *S. aureus* at 40 labels, 0.008, 0.003 to 0.013). Since the table shows that uncertainty sampling lost AUROC in more than one configuration, the sentence in Section VI-B that called it no worse than random selection in all but one configuration now gives the average cost instead (0.017 for *K. pneumoniae*, 0.006 for *S. aureus* and nothing measurable for *E. coli*).

For the cross site setting the text no longer says that no comparison reached significance, because with 5 seeds none could have. Those comparisons are now read as effect sizes (Section V-B): the mean difference was negative in 64 of the 90 informative comparisons, and its interval excluded zero in 7 (6 against the rule and 1 for it). All 135 comparisons, with their intervals, raw p values and both corrections, are in the released file `paper/camera_ready/stats`, together with the script that computes them.

2. Demonstrate that the negative findings are robust to the hash based sequential batching.

A new experiment was run for this point (Section V-B, last paragraph). The temporal comparison at 100, 400 and 800 labels was repeated with the same data, model, tuned settings and seeds, and only the division of the 2017 pool changed. Three further divisions were drawn from independent hashes of the isolate code, and in a fourth setting the sequence was removed and the whole budget was bought in a single round. The division of the main experiment was rerun alongside as a reference and reproduced all 360 of the corresponding published runs exactly, which also confirms that the simulation is reproducible.

The negative finding held in every division. The mean difference against random selection was negative in 16 to 20 of the 27 comparisons of each division. After correction within each pair, 8 comparisons across

the 5 divisions favoured random selection and 1 favoured a rule (diversity sampling on *K. pneumoniae* at 800 labels in one redrawn division, better by 0.025, 95% confidence interval 0.014 to 0.036). We report that comparison openly in the paper. It was positive in all 5 divisions, by 0.003 to 0.025, and is the only advantage of any rule that recurred. The safety result was equally stable, since uncertainty sampling at 800 labels reduced the resistant isolates reported susceptible in every division and for every pair, with an interval excluding zero in all 15 cases.

We also state plainly what this experiment cannot show. The released spectra carry no acquisition dates, so no division reproduces a calendar order (Section VI-C).

3. Provide sufficient cross hospital experimental details and reproducible simulation accounting.

Section IV-A now gives the cross site protocol: the source model is fitted on the DRIAMS-A isolates of 2015 to 2018, the 2018 isolates of each target site are divided in half (stratified by label and redrawn for every seed) into a pool and a test set, the budget (25 to 400 labels) is spent in a single round, and the model is refitted on the source isolates together with the bought ones. *S. aureus* with oxacillin has no labelled isolates at DRIAMS-D, which leaves 8 combinations of pair and site.

Preparing these details showed something the submitted version did not disclose. In 30 of the 120 cross site comparisons the budget was at least as large as the pool of the target site, so every rule bought the same isolates and the difference is zero by construction. The informative comparisons are therefore 90, and 135 in all rather than 165. The abstract, the introduction, the results and the conclusion now use 135.

Section IV-B now splits the 1540 simulations into 660 temporal and 880 cross site. In full, these are 660 temporal, from 3 pairs, 22 configurations made of 4 rules at 5 budgets and the 2 bounds, and 10 seeds, and 880 cross site, from 8 combinations of pair and site, 22 configurations and 5 seeds. Three descriptions in the method were also corrected so that they match the code that produced the results: the hybrid rule (the $4b$ isolates closest to the threshold, grouped into b clusters by k-means, one bought per cluster), the diversity rule (k-center on the first 50 principal components) and the threshold (the highest value that still gives 90% sensitivity). No result changed because of these corrections.

Reviewer 1

1. Clarify the limitations of the simulated data ordering.

Section IV-A now states how the batches are formed (the MD5 digest of the isolate code modulo 4) and that their order carries no calendar information. Section VI-C now states that the simulation represents drift between years but not within a year, and that a rule which adapts to recent isolates could gain from a real ordering in a way these experiments cannot show. Redrawing the batches (Section V-B) establishes that the result does not depend on one particular division. The words “quarterly rounds” were replaced by “rounds” in the abstract and the introduction.

2. Provide more details about the statistical comparisons.

See the reply to Reviewer 2, point 1. Section IV-C gives the test, the pairing, the number of seeds, both correction families and the smallest attainable p values, and Table I gives the mean paired difference with its 95% confidence interval for every temporal comparison.

3. Discuss the practical implications of the increased major errors for some configurations.

Section VI-B now states what the trade means in practice. For *K. pneumoniae* at 800 labels it is about 5 additional major errors for each very major error avoided (66 against 14), and since a major error lasts only until the culture based result arrives, a laboratory may accept it, but it is a trade rather than a free gain. The same paragraph reports that on the 2018 isolates the models maintained by uncertainty sampling made more very major errors and fewer major errors than those maintained by random selection in 14 of the 15 configurations, and recommends that a laboratory using the rule set the threshold again on recent labels.

4. Discuss the need for validation on real time clinical data.

Section VI-C now states that the study is a retrospective simulation in which a bought label is available at once, and that the safety result needs confirmation on isolates collected in real time before it is used to set a testing policy. The conclusion names validation on isolates with real acquisition dates as future work.

Editor

Format. The paper is now in the IEEE two column conference format on A4. The author names, affiliations and emails are given without titles. The notice 979-8-3195-2818-6/26/$31.00 ©2026 IEEE is at the bottom left of the first page only. The title and the author list are unchanged, and the reference list is unchanged (20 references).

Length. The paper is 6 pages in the IEEE A4 format.

Proofreading. The whole paper was proofread again.

Use of AI tools. As required, the use of an AI assistant is stated in the Acknowledgment.